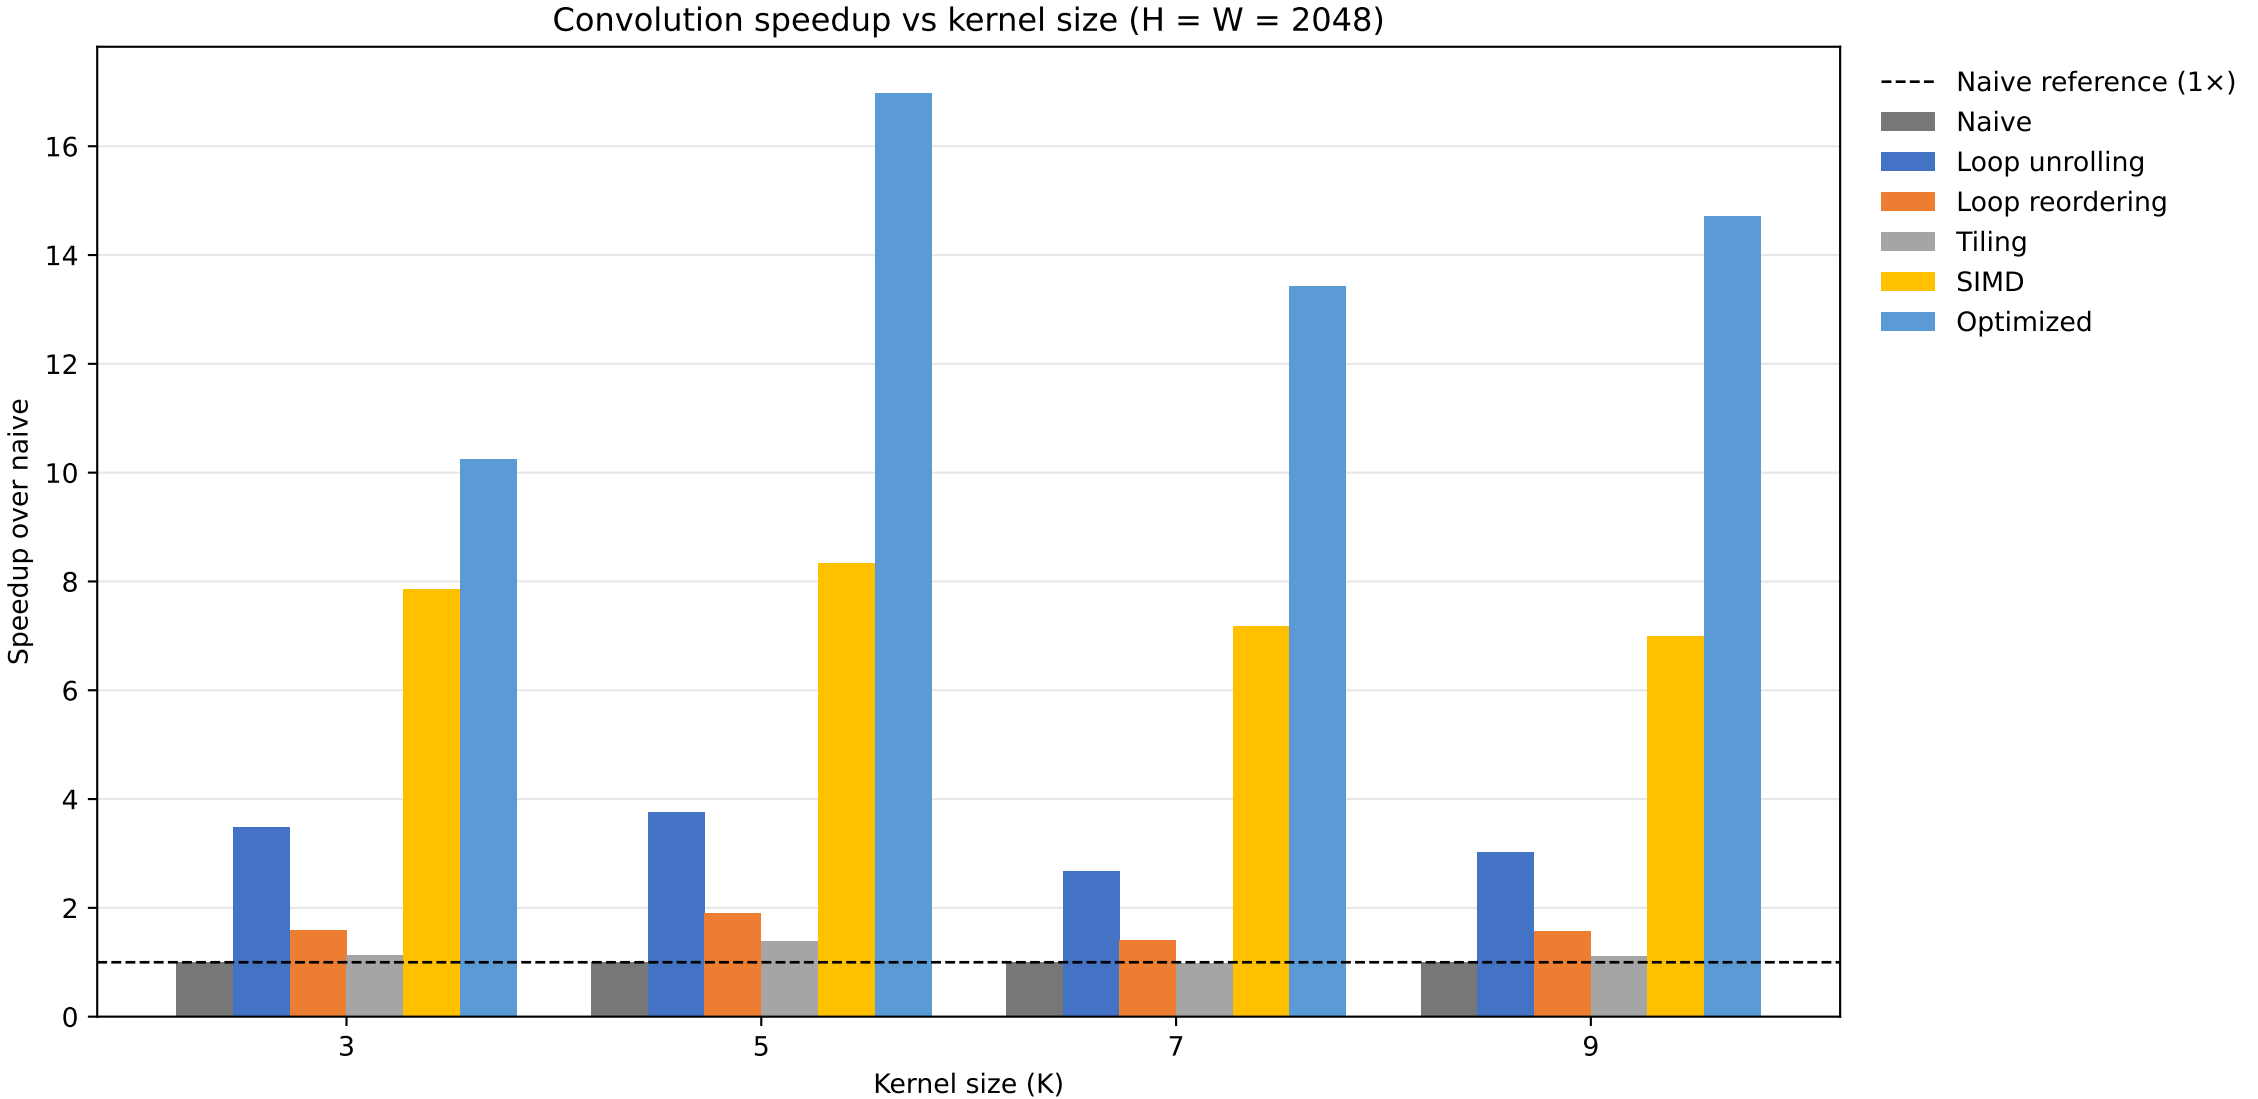
Convolution speedup vs kernel size (H = W = 2048)
Speedup over naive
Kernel size (K)
0
2
4
6
8
10
12
14
16
3
5
7
9
Naive reference (1×)
Naive
Loop unrolling
Loop reordering
Tiling
SIMD
Optimized